

ACKNOWLEDGMENTS

This work was supported by NSF grant IIS-2239130.

REFERENCES

- [1] Google Images, the most comprehensive image search on the web., 2023. https://www.google.com/advanced_image_search. 2
- [2] Plotly: the front end for ML and data science models, 2023. <https://plotly.com/>. 3, 7
- [3] See it, search it — Bing Visual Search, 2023. <https://www.bing.com/visualsearch>. 2
- [4] AnyChart. Anychart is a lightweight and robust javascript charting library, 2025. <https://www.anychart.com/>. 6
- [5] ApexCharts. Apexcharts.js - open source javascript charts for your website, 2024. <https://apexcharts.com/>. 3
- [6] ApexCharts. Angular chart examples & samples demo, 2025. <https://apexcharts.com/angular-chart-demos/>. 6
- [7] ApexCharts.js. A Stacked Bar Chart. <https://apexcharts.com/javascript-chart-demos/column-charts/stacked/>. 8
- [8] L. Battle, P. Duan, Z. Miranda, D. Mukusheva, R. Chang, and M. Stonebraker. Beagle: Automated Extraction and Interpretation of Visualizations from the Web. In *Proceedings of the 2018 CHI Conference on Human Factors in Computing Systems*, article no. 594, pp. 594:1–594:8. ACM, Montreal, 2018. doi: 10.1145/3173574.3174168 2, 3, 5
- [9] M. A. Borkin, A. A. Vo, Z. Bylinskii, P. Isola, S. Sunkavalli, A. Oliva, and H. Pfister. What makes a visualization memorable? *IEEE Transactions on Visualization and Computer Graphics*, 19(12):2306–2315, 2013. doi: 10.1109/TVCG.2013.234 5
- [10] M. Bostock. bl.ocks.org, 2021. <https://bl.ocks.org/>. 2
- [11] M. Bostock, V. Ogievetsky, and J. Heer. D³ Data-Driven Documents. *IEEE Transactions on Visualization and Computer Graphics*, 17(12):2301–2309, 2011. doi: 10.1109/TVCG.2011.185 2, 3, 6
- [12] V. S. Bursztyn, J. Hoffswell, E. Koh, and S. Guo. Representing charts as text for language models: An in-depth study of question answering for bar charts. In *2024 IEEE Visualization and Visual Analytics (VIS)*, pp. 266–270, 2024. doi: 10.1109/VIS55277.2024.00061 6
- [13] M. Caron, H. Touvron, I. Misra, H. Jegou, J. Mairal, P. Bojanowski, and A. Joulin. Emerging properties in self-supervised vision transformers. In *2021 IEEE/CVF International Conference on Computer Vision (ICCV)*, pp. 9630–9640, 2021. doi: 10.1109/ICCV48922.2021.00951 8
- [14] Chartmaker. The chartmaker directory, 2021. <http://chartmaker.visualisingdata.com/>. 2
- [15] C. Chen, B. Lee, Y. Wang, Y. Chang, and Z. Liu. Mystique: Deconstructing svg charts for layout reuse. *IEEE Transactions on Visualization and Computer Graphics*, 30(1):447–457, 2024. doi: 10.1109/TVCG.2023.3327354 1, 2, 4, 5, 6, 7
- [16] C. Chen and Z. Liu. The State of the Art in Creating Visualization Corpora for Automated Chart Analysis. *Computer Graphics Forum*, 42(3):449–470, 2023. doi: 10.1111/cgf.14855 1, 2, 3, 4, 5, 7, 9
- [17] J. Chen, M. Ling, R. Li, P. Isenberg, T. Isenberg, M. Sedlmair, T. Möller, R. S. Laramée, H.-W. Shen, K. Wünsche, and Q. Wang. Vis30k: A collection of figures and tables from ieee visualization conference publications. *IEEE Transactions on Visualization and Computer Graphics*, 27(9):3826–3833, 2021. doi: 10.1109/TVCG.2021.3054916 5
- [18] N. Chen, Y. Zhang, J. Xu, K. Ren, and Y. Yang. Viseval: A benchmark for data visualization in the era of large language models. *IEEE Transactions on Visualization and Computer Graphics*, 31(1):1301–1311, 11 pages, Jan. 2025. doi: 10.1109/TVCG.2024.3456320 5
- [19] X. Chen, W. Zeng, Y. Lin, H. M. Al-manee, J. Roberts, and R. Chang. Composition and configuration patterns in multiple-view visualizations. *IEEE Transactions on Visualization and Computer Graphics*, 27(2):1514–1524, 2021. doi: 10.1109/TVCG.2020.3030338 5
- [20] C. Cortes and V. Vapnik. Support-vector networks. *Mach. Learn.*, 20(3):273–297, 25 pages, Sept. 1995. doi: 10.1023/A:1022627411411 7
- [21] W. Cui, J. Wang, H. Huang, Y. Wang, C.-Y. Lin, H. Zhang, and D. Zhang. A Mixed-Initiative Approach to Reusing Infographic Charts. *IEEE Transactions on Visualization and Computer Graphics*, 28(1):173–183, 2022. doi: 10.1109/TVCG.2021.3114856 1, 7
- [22] W. Cui, X. Zhang, Y. Wang, H. Huang, B. Chen, L. Fang, H. Zhang, J.-G. Lou, and D. Zhang. Text-to-viz: Automatic generation of infographics from proportion-related natural language statements. *IEEE Transactions on Visualization and Computer Graphics*, 26(1):906–916, 2019. doi: 10.1109/TVCG.2019.2934785 1
- [23] D3.js. A Geographical Heatmap. <https://observablehq.com/@z-richard/choropleth-map-of-the-covid-19-cases-in-china>. 8
- [24] Datylon. Datylon: Design, automate & share beautiful, on-brand reports, 2024. <https://www.datylon.com/>. 3
- [25] K. Davila, S. Setlur, D. Doermann, B. U. Kota, and V. Govindaraju. Chart mining: A survey of methods for automated chart analysis. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 43(11):3799–3819, 2020. doi: 10.1109/TPAMI.2020.2992028 1
- [26] D. Davis. An Annotated Bar Chart. <https://poetrybetweenpain.deb.is/>. 8
- [27] D. Deng, Y. Wu, X. Shu, J. Wu, S. Fu, W. Cui, and Y. Wu. VisImages: A Fine-Grained Expert-Annotated Visualization Dataset. *IEEE Transactions on Visualization & Computer Graphics*, 29(07):3298–3311, July 2023. doi: 10.1109/TVCG.2022.3155440 2, 5, 7
- [28] V. Dibia and Ç. Demiralp. Data2vis: Automatic generation of data visualizations using sequence-to-sequence recurrent neural networks. *IEEE Computer Graphics and Applications*, 39(5):33–46, 2019. doi: 10.1109/MCG.2019.2924636 1
- [29] S. Dou, X. Jiang, L. Liu, L. Ying, C. Shan, Y. Shen, X. Dong, Y. Wang, D. Li, and C. Zhao. Hierarchically recognizing vector graphics and a new chart-based vector graphics dataset. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 46(12):7556–7573, 2024. doi: 10.1109/TPAMI.2024.3394298 5
- [30] Ferdio. The data viz project, 2021. <https://datavizproject.com/>. 2
- [31] J. Harper and M. Agrawala. Deconstructing and restyling D3 visualizations. In *Proceedings of the 27th annual ACM symposium on User interface software and technology*, pp. 253–262. ACM, Honolulu, 2014. doi: 10.1145/2642918.2647411 7
- [32] J. Heer, M. Agrawala, and W. Willett. Generalized selection via interactive query relaxation. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*, CHI ’08, 10 pages, p. 959–968. Association for Computing Machinery, New York, NY, USA, 2008. doi: 10.1145/1357054.1357203 4
- [33] Highcharts. A Multi-Line Chart. <https://www.highcharts.com/demo/highcharts/line-ajax>. 8
- [34] Highcharts. Highcharts: Interactive charting library, 2024. <https://www.highcharts.com/>. 3
- [35] A. Howard, M. Sandler, B. Chen, W. Wang, L.-C. Chen, M. Tan, G. Chu, V. Vasudevan, Y. Zhu, R. Pang, H. Adam, and Q. Le. Searching for MobileNetV3. In *2019 IEEE/CVF International Conference on Computer Vision (ICCV)*, pp. 1314–1324. IEEE Computer Society, Los Alamitos, CA, USA, Nov. 2019. doi: 10.1109/ICCV.2019.00140 8
- [36] K. Hu, M. A. Bakker, S. Li, T. Kraska, and C. Hidalgo. Vizml: A Machine Learning Approach to Visualization Recommendation. In *Proceedings of the 2019 CHI Conference on Human Factors in Computing Systems*, pp. 1–12, 2019. doi: 10.1145/3290605.3300358 1
- [37] K. Hu, S. N. S. Gaikwad, M. Hulsebos, M. A. Bakker, E. Zraggen, C. Hidalgo, T. Kraska, G. Li, A. Satyanarayan, and C. Demiralp. VizNet: Towards A Large-Scale Visualization Learning and Benchmarking Repository. In *Proceedings of the 2019 CHI Conference on Human Factors in Computing Systems*, pp. 1–12. ACM, Glasgow Scotland Uk, May 2019. doi: 10.1145/3290605.3300892 9
- [38] A. Hurst, A. Lerer, A. P. Goucher, A. Perelman, A. Ramesh, A. Clark, A. Ostrow, A. Welihinda, A. Hayes, A. Radford, et al. GPT-4o System Card. *arXiv preprint arXiv:2410.21276*, 2024. 6
- [39] D. Jung, W. Kim, H. Song, J.-I. Hwang, B. Lee, B. Kim, and J. Seo. Chartsense: Interactive Data Extraction from Chart Images. In *Proceedings of the 2017 CHI Conference on Human Factors in Computing Systems*, 12 pages, pp. 6706–6717. ACM, Denver, 2017. doi: 10.1145/3025453.3025957 1, 7
- [40] S. E. Kahou, V. Michalski, A. Atkinson, Á. Kádár, A. Trischler, and Y. Bengio. Figureqa: An annotated figure dataset for visual reasoning. *arXiv preprint arXiv:1710.07300*, 2017. 9
- [41] G. Kazai, J. Kamps, and N. Milic-Frayling. An analysis of human factors and label accuracy in crowdsourcing relevance judgments. *Inf. Retr.*, 16(2):138–178, 41 pages, Apr. 2013. doi: 10.1007/s10791-012-9205-0 3
- [42] D. P. Kingma and J. Ba. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*, 2014. 8
- [43] A. Kittur, E. H. Chi, and B. Suh. Crowdsourcing user studies with me-

- chanical turk. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*, CHI '08, 4 pages, p. 453–456. Association for Computing Machinery, New York, NY, USA, 2008. doi: [10.1145/1357054.1357127](https://doi.org/10.1145/1357054.1357127) 3
- [44] H.-K. Ko, H. Jeon, G. Park, D. H. Kim, N. W. Kim, J. Kim, and J. Seo. Natural language dataset generation framework for visualizations powered by large language models. In *Proceedings of the CHI Conference on Human Factors in Computing Systems*, CHI '24, article no. 843, 22 pages. Association for Computing Machinery, New York, NY, USA, 2024. doi: [10.1145/3613904.3642943](https://doi.org/10.1145/3613904.3642943) 5
- [45] Y. LeCun and Y. Bengio. *Convolutional networks for images, speech, and time series*, p. 255–258. MIT Press, Cambridge, MA, USA, 1998. 7
- [46] H. Li, Y. Wang, A. Wu, H. Wei, and H. Qu. Structure-aware visualization retrieval. In *Proceedings of the 2022 CHI Conference on Human Factors in Computing Systems*, CHI '22, article no. 409, 14 pages. Association for Computing Machinery, New York, NY, USA, 2022. doi: [10.1145/3491102.3502048](https://doi.org/10.1145/3491102.3502048) 1, 2, 7
- [47] A. Liu, B. Feng, B. Xue, B. Wang, B. Wu, C. Lu, C. Zhao, C. Deng, C. Zhang, C. Ruan, et al. DeepSeek-V3 Technical Report. *arXiv preprint arXiv:2412.19437*, 2024. 6
- [48] Z. Liu, C. Chen, and J. Hooker. Manipulable Semantic Components: A Computational Representation of Data Visualization Scenes. *IEEE Transactions on Visualization and Computer Graphics*, 31(1):732 – 742, 2025. doi: [10.1109/TVCG.2024.3456296](https://doi.org/10.1109/TVCG.2024.3456296) 2, 3, 4, 6, 9
- [49] Z. Liu, C. Chen, F. Morales, and Y. Zhao. Atlas: Grammar-based Procedural Generation of Data Visualizations. In *2021 IEEE Visualization Conference (VIS)*, pp. 171–175. IEEE, New Orleans, 2021. doi: [10.1109/VIS49827.2021.9623315](https://doi.org/10.1109/VIS49827.2021.9623315) 2, 3
- [50] S. Madan, Z. Bylinskii, M. Tancik, A. Recasens, K. Zhong, S. Alsheikh, H. Pfister, A. Oliva, and F. Durand. Synthetically trained icon proposals for parsing and summarizing infographics. *arXiv preprint arXiv:1807.10441*, 2018. 5
- [51] A. Mahinpei, Z. Kostic, and C. Tanner. Linecap: Line charts for data visualization captioning models. In *2022 IEEE Visualization and Visual Analytics (VIS)*, pp. 35–39. IEEE, 2022. doi: [10.1109/VIS54862.2022.00016](https://doi.org/10.1109/VIS54862.2022.00016) 9
- [52] Mascot.js. A Connected Dot Plot. <https://mascot-vis.github.io/gallery#DumbbellChart>. 8
- [53] D. Masson, S. Malacria, D. Vogel, E. Lank, and G. Casiez. ChartDetective: Easy and Accurate Interactive Data Extraction from Complex Vector Charts. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems*, number 147, article no. 147, 17 pages. ACM, Hamburg, 2023. doi: [10.1145/3544548.3581113](https://doi.org/10.1145/3544548.3581113) 2, 9
- [54] Matplotlib. Matplotlib — visualization with python, 2025. <https://matplotlib.org/>. 6
- [55] A. M. McNutt. No Grammar to Rule Them All: A Survey of JSON-style DSLs for Visualization. *IEEE Transactions on Visualization and Computer Graphics*, 29(1):160–170, 2023. doi: [10.1109/TVCG.2022.3209460](https://doi.org/10.1109/TVCG.2022.3209460) 2
- [56] Microsoft. Azure AI Vision with OCR and AI - Microsoft Azure. <https://azure.microsoft.com/en-us/products/ai-services/ai-vision>. 2
- [57] S. Minaee, Y. Boykov, F. Porikli, A. Plaza, N. Kehtarnavaz, and D. Terzopoulos. Image segmentation using deep learning: A survey. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 44(7):3523–3542, 2022. doi: [10.1109/TPAMI.2021.3059968](https://doi.org/10.1109/TPAMI.2021.3059968) 2
- [58] T. Munzner. *Visualization analysis and design*. CRC press, 2014. 2
- [59] NIVO. Home — nivo, 2024. <https://nivo.rocks/>. 3, 6
- [60] Observable. Explore, analyze, and explain data as a team., 2021. <https://observablehq.com/>. 2
- [61] J. Poco and J. Heer. Reverse-Engineering Visualizations: Recovering Visual Encodings from Chart Images. *Computer Graphics Forum*, 36(3):353–363, 2017. doi: [10.1111/cgf.13193](https://doi.org/10.1111/cgf.13193) 1, 2, 5, 6, 7
- [62] D. Ren, B. Lee, and M. Brehmer. Charticulator: Interactive Construction of Bespoke Chart Layouts. *IEEE Transactions on Visualization and Computer Graphics*, 25(1):789–799, 2018. doi: [10.1109/TVCG.2018.2865158](https://doi.org/10.1109/TVCG.2018.2865158) 4, 9
- [63] S. Ribeca. The data visualisation catalogue, 2021. <https://datavizcatalogue.com/about.html>. 2
- [64] A. Satyanarayan, B. Lee, D. Ren, J. Heer, J. Stasko, J. Thompson, M. Brehmer, and Z. Liu. Critical Reflections on Visualization Authoring Systems. *IEEE Transactions on Visualization and Computer Graphics*, 26(1):461–471, 2019. doi: [10.1109/TVCG.2019.2934281](https://doi.org/10.1109/TVCG.2019.2934281) 2, 3
- [65] A. Satyanarayan, D. Moritz, K. Wongsuphasawat, and J. Heer. Vega-Lite: A Grammar of Interactive Graphics. *IEEE Transactions on Visualization and Computer Graphics*, 23(1):341–350, 2016. doi: [10.1109/TVCG.2016.2599030](https://doi.org/10.1109/TVCG.2016.2599030) 2, 3, 9
- [66] M. Savva, N. Kong, A. Chhajta, L. Fei-Fei, M. Agrawala, and J. Heer. Revision: Automated Classification, Analysis and Redesign of Chart Images. In *Proceedings of the 24th annual ACM symposium on User interface software and technology*, pp. 393–402. ACM, Santa Barbara, 2011. doi: [10.1145/2047196.2047247](https://doi.org/10.1145/2047196.2047247) 1, 2, 6, 7
- [67] F. Scarselli, M. Gori, A. C. Tsoi, M. Hagenbuchner, and G. Monfardini. The graph neural network model. *IEEE Transactions on Neural Networks*, 20(1):61–80, 2009. doi: [10.1109/TNN.2008.2005605](https://doi.org/10.1109/TNN.2008.2005605) 7
- [68] Semiotic. Semiotic — a data visualization framework for react, 2025. <https://semiotic.interact.io/>. 6
- [69] L. S. Snyder and J. Heer. DIVI: Dynamically Interactive Visualization. *IEEE Transactions on Visualization and Computer Graphics*, 30(1):403 – 413, 2024. Publisher: IEEE. doi: [10.1109/TVCG.2023.3327172](https://doi.org/10.1109/TVCG.2023.3327172) 2, 3
- [70] Spotfire. Spotfire: Transforming Data into Real-Time Insights and Actionable Decisions. <https://www.spotfire.com/>. 2
- [71] B. J. Tang, A. Boggust, and A. Satyanarayan. VisText: A Benchmark for Semantically Rich Chart Captioning. In *The Annual Meeting of the Association for Computational Linguistics (ACL)*, 2023. 5
- [72] M. USA. Award-winning fast, flexible c# .net chart control, 2025. <https://developer.mescius.com/componentone/docs/win/online-flexchart/overview.html>. 6
- [73] W3C (2019). WAI Web Accessibility Tutorials: Complex Images., 2019. <https://www.w3.org/WAI/tutorials/images/complex/>. 8
- [74] Z. Wang, J. Hsu, X. Wang, K.-H. Huang, M. Li, J. Wu, and H. Ji. Text-based reasoning about vector graphics. *arXiv preprint arXiv:2404.06479*, 2024. 6
- [75] L. Wilkinson. ggplot2: elegant graphics for data analysis by wickham, h., 2011. 3
- [76] K. Xu, W. Hu, J. Leskovec, and S. Jegelka. How Powerful Are Graph Neural Networks? In *International Conference on Learning Representations*, 2019. 6, 7
- [77] S. Yun, M. Jeong, R. Kim, J. Kang, and H. J. Kim. Graph Transformer Networks. *Advances in Neural Information Processing Systems*, 32, 2019. 8
- [78] K. T. Zinat, J. Yang, A. Gandhi, N. Mitra, and Z. Liu. A comparative evaluation of visual summarization techniques for event sequences. *Computer Graphics Forum*, 42(3):173–185, 2023. doi: [10.1111/cgf.14821](https://doi.org/10.1111/cgf.14821) 3
- [79] J. Zong, C. Lee, A. Lundgard, J. Jang, D. Hajas, and A. Satyanarayan. Rich Screen Reader Experiences for Accessible Data Visualization. *Computer Graphics Forum (Proc. EuroVis)*, 2022. doi: [10.1111/cgf.14519](https://doi.org/10.1111/cgf.14519) 1, 6, 8